

Fair Exponential Smoothing with Small Alpha

(fessa)

Juha Reunanen
juha.reunanen@tomaattinen.com

October 10, 2015

Abstract

This short paper shows that standard exponential smoothing with a small α may not give the desired results $s(t)$ for small values of t , because the first sample $x(0)$ has such a huge impact on the outcome. Fortunately, the paper also proposes a remedy to the problem.

1 Introduction

Exponential smoothing¹ is a simple and effective technique for low-pass filtering a time series. In its most basic form, it is defined like this:

$$\begin{aligned}s(0) &= x(0) \\ s(t) &= \alpha \cdot x(t) + (1 - \alpha) \cdot s(t - 1)\end{aligned}$$

where $x(t)$ is the raw input signal, $s(t)$ is the filtered signal, and α is a parameter ($0 < \alpha < 1$).

The technique has many nice properties, one of which is that even with a small α (i.e., heavy filtering), finding $s(t)$ is very fast — compared to calculating a moving average using a long window. This is very nice in case $x(t)$ is a sequence of large images, for instance.

2 Problem

One problem with using exponential smoothing, in particular with a small α , is that if $x(0)$ happens to be an outlier, it affects $s(t)$ even for pretty large values of t .

Denote with $w(k, t)$ the effect that $x(k)$ has over $s(t)$. Clearly:

$$\begin{aligned}w(0, 0) &= 1 \\ w(0, 1) &= 1 - \alpha \\ w(1, 1) &= \alpha \\ w(0, 2) &= (1 - \alpha)^2 \\ w(1, 2) &= \alpha(1 - \alpha) \\ w(2, 2) &= \alpha^2 \\ &\dots\end{aligned}$$

Note that $w(k, t) = 0$ for any $k > t$, because we have a causal filter (i.e., events from the future cannot affect the current filtered value).

Now we can see the problem: the effect of the first sample, $w(0, t) = (1 - \alpha)^t$, is quite significant even for large t , if α is small.

¹https://en.wikipedia.org/wiki/Exponential_smoothing

Consider for instance $\alpha = 0.001$. You need to wait until $t = 693$ before the relative weight of the first sample drops below 50%! So the 691 samples $x(1), \dots, x(691)$, all newer than $x(0)$, together weigh less than $x(0)$ does alone! And what if you want to use an even smaller α ? It does not appear far-fetched to say that this can be quite unfair, depending of course on what exactly it is that you are doing.

3 Challenge

So we want to define a new, fairer scheme for $w(k, t)$. What properties should we expect?

Clearly we need to have $\sum_{k=0}^t w(k, t) = 1$ for any $t \geq 0$. Let's call this our first requirement.

Also, we would like to keep this asymptotic property of standard exponential smoothing:

$$w(k-1, t) = (1 - \alpha) \cdot w(k, t)$$

But why not just require it for any k and t ? Standard exponential smoothing caters for this only when $k > 2$ (any $t \geq k$) and $t \rightarrow \infty$ (any k). Yes, let's try and call this our second requirement.

Finally, we would like to be able to compute a new $s(t)$ using a small number of operations only.

4 Solution

Putting together the first and the second requirement, we have:

$$\left\{ \begin{array}{l} \sum_{k=0}^t w(k, t) = 1 \\ w(k-1, t) = (1 - \alpha) \cdot w(k, t) \end{array} \right. \quad (1)$$

Solving, it appears that we get:

$$w(k, t) = \frac{(1 - \alpha)^{t-k}}{\sum_{\tau=0}^t (1 - \alpha)^\tau}$$

This holds for $k \leq t$. For $k > t$, we have $w(k, t) = 0$, as pointed out earlier.

For the `fessa` method being introduced², we are looking for an $\alpha_{\text{fessa}}(t)$ that we can use exactly as we use α in the standard method. Now we can define $\alpha_{\text{fessa}}(t)$ as follows:

$$\alpha_{\text{fessa}}(t) = w(t, t) = \frac{1}{\sum_{\tau=0}^t (1 - \alpha)^\tau} = \frac{1}{d(t)}$$

Turns out that we can update α_{fessa} recursively:

$$\begin{aligned} a(0) &= d(0) = \alpha_{\text{fessa}}(0) = 1 \\ a(t) &= a(t-1) \cdot (1 - \alpha) \\ d(t) &= d(t-1) + a(t) \\ \alpha_{\text{fessa}}(t) &= \frac{1}{d(t)} \end{aligned}$$

The filtering itself works as previously — we just use $\alpha_{\text{fessa}}(t)$ instead of α :

$$s(t) = \alpha_{\text{fessa}}(t) \cdot x(t) + (1 - \alpha_{\text{fessa}}(t)) \cdot s(t-1)$$

²`fessa` stands for *fair exponential smoothing with small alpha*.

Table 1: Table of effective $w(k, t)$ in standard exponential smoothing for different k and $t \leq 10$, when $\alpha = 0.001$. The numbers are rounded to four significant digits. It is clear how $w(0, t)$ dominates — which means that the value of $x(0)$ has a significant impact on $s(t)$ when t is small.

	t											
	0	1	2	3	4	5	6	7	8	9	10	
k	0	1.000	.9990	.9980	.9970	.9960	.9950	.9940	.9930	.9920	.9910	.9900
	1		.0010	.0010	.0010	.0010	.0010	.0010	.0010	.0010	.0010	.0010
	2			.0010	.0010	.0010	.0010	.0010	.0010	.0010	.0010	.0010
	3				.0010	.0010	.0010	.0010	.0010	.0010	.0010	.0010
	4					.0010	.0010	.0010	.0010	.0010	.0010	.0010
	5						.0010	.0010	.0010	.0010	.0010	.0010
	6							.0010	.0010	.0010	.0010	.0010
	7								.0010	.0010	.0010	.0010
	8									.0010	.0010	.0010
	9										.0010	.0010
	10											.0010

Table 2: Table of effective $w(k, t)$ in the `fessa` method for different k and $t \leq 10$, when $\alpha = 0.001$. Compared to Table 1, we can see that for each t , all samples received so far have an almost equal significance, however a more recent sample always weighs slightly more than do the earlier values of x .

	t											
	0	1	2	3	4	5	6	7	8	9	10	
k	0	1.000	.4997	.3330	.2496	.1996	.1663	.1424	.1246	.1107	.0996	.0905
	1		.5003	.3333	.2499	.1998	.1664	.1426	.1247	.1108	.0997	.0905
	2			.3337	.2501	.2000	.1666	.1427	.1248	.1109	.0997	.0906
	3				.2504	.2002	.1667	.1429	.1249	.1110	.0998	.0907
	4					.2004	.1669	.1430	.1251	.1111	.0999	.0908
	5						.1671	.1431	.1252	.1112	.1000	.0909
	6							.1433	.1253	.1113	.1001	.0910
	7								.1254	.1114	.1003	.0911
	8									.1116	.1004	.0912
	9										.1005	.0913
	10											.0914

5 Examples

For $\alpha = 0.001$, Tables 1 and 2 show the effective $w(k, t)$ for small values of k and t for both the standard method and the proposed `fessa` approach, respectively. Based on Table 1, it is clear that in standard exponential smoothing the value of $x(0)$ has a huge impact on $s(t)$, when t is small. On the other hand, based on Table 2, the properties of the `fessa` method appear to be much nicer: at any time t , all samples processed so far have roughly the same weight, however the more recent the sample, the (slightly) more it weighs.

As another simple illustration, consider $x(t)$ drawn from a uniform distribution between 0 and 100. Have a look at Figures 1–3. It is clear that with such a small alpha (0.001), whatever $x(0)$ happens to be affects the results of the standard exponential smoothing method for a very long time — too long, one could probably say. On the other hand, the `fessa` method being proposed finds its way near the expected value quite soon.

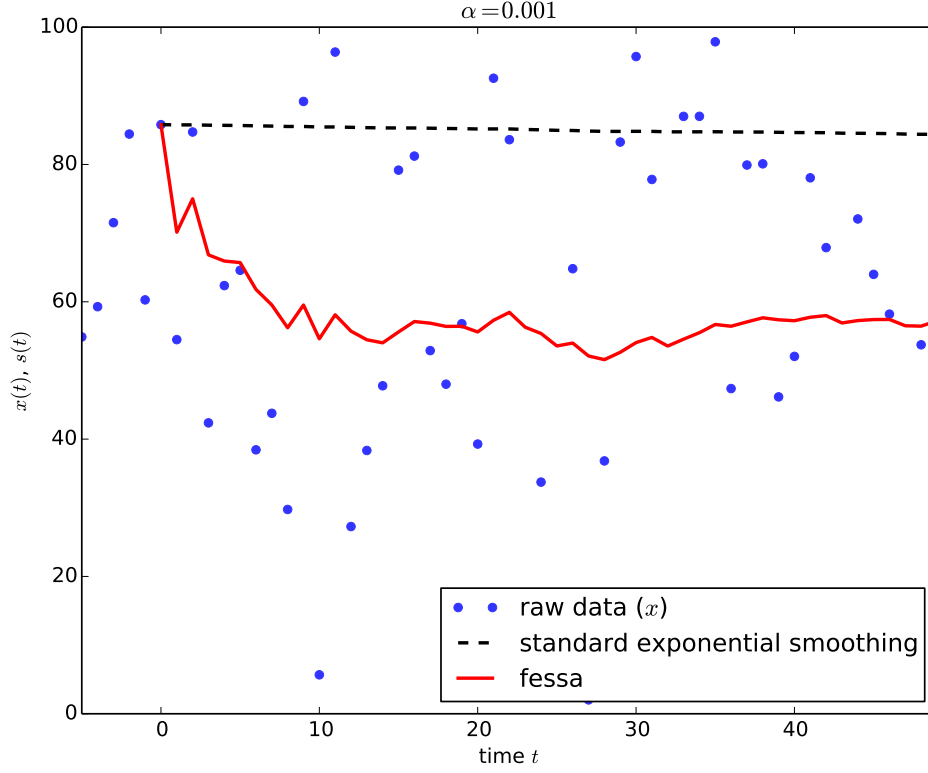


Figure 1: Comparison between standard exponential smoothing and the proposed fessa method for small t . Note that the standard method gets seriously stuck with the $x(0)$, even though it is not really even an outlier.




Figure 2: Comparison between standard exponential smoothing and the proposed fessa method for a medium range of t . Given that the expected value is 50, the standard method is still far away, even though the direction is clearly right.

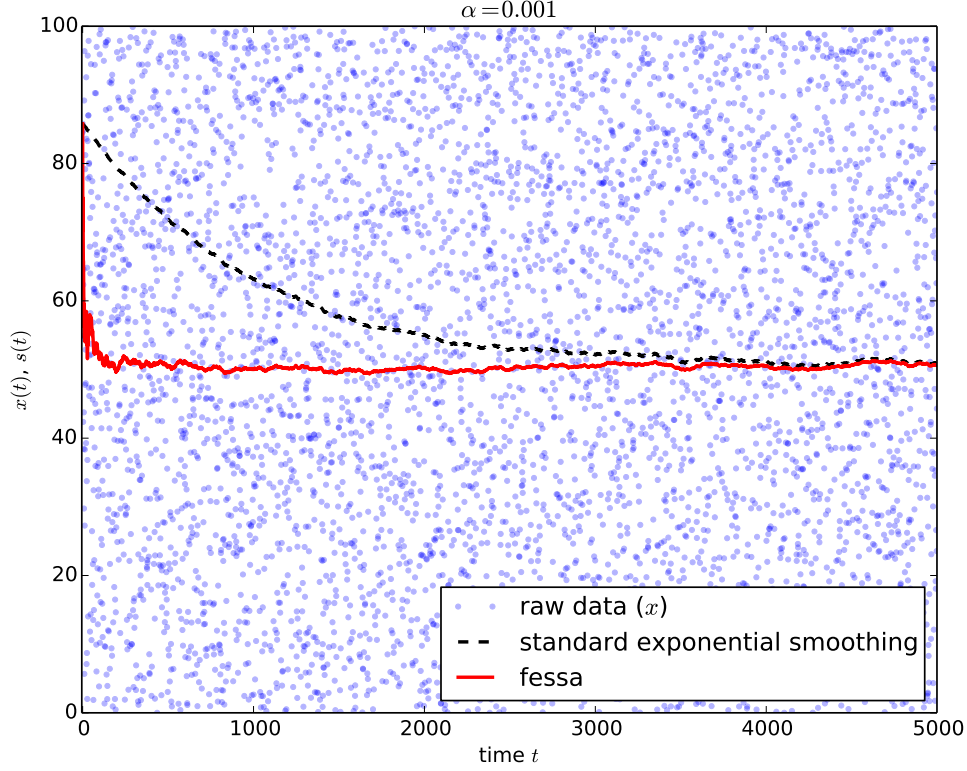


Figure 3: Comparison between standard exponential smoothing and the proposed fessa method for large t . Now even the standard method has reached the expected value. From around here on, the two methods are essentially identical.

6 Runtime performance

Because asymptotically ($t \rightarrow \infty$) the fessa method is equivalent to standard exponential smoothing, any implementation can switch to the standard method as soon as $(1 - \alpha)^t$ gets really small (close to machine epsilon or so). Then, for a large t the computational cost is on the same level with standard exponential smoothing.

With a small t , compared to the standard method there is additional multiplication (to update a), addition (to update d) and division³ (to calculate α_{fessa}), so do not use fessa if the overhead is too much for you.⁴ Otherwise, feel free to use the proposed method in place of standard exponential smoothing.

7 Missing proofs

It would be nice to prove that $\lim_{t \rightarrow \infty} \alpha_{\text{fessa}}(t) = \alpha$.

It is quite obvious that $w(k, t) = \frac{(1-\alpha)^{t-k}}{\sum_{\tau=0}^t (1-\alpha)^\tau}$ satisfies (1), but should it be shown that the solution is unique?

³Actually this is just the reciprocal, which may be faster to calculate than generic division.

⁴But in this case please let me know what it is that you are working with — it sounds quite interesting already!

8 Disclaimer

The technique presented in this paper may well be a well-known one, but nevertheless I had fun figuring it out — not to mention coming up with the name. Indeed I think I need the method in practice, since I have a few times hesitated to set a low α because of being afraid of a very noisy $x(0)$. So if after this we can set lower α s and get smoother signals, it is quite enough for me, regardless of whether someone came up with this method already in the 1600s or so.

9 Code

See <https://github.com/reunanen/fessa> for a simple implementation of the `fessa` technique.